AI AGENTS · CAREER TOOLS

A free AI that tailors resumes, and code that checks it.

Real runs on Brex and BILL postings, mistakes included.

· Julian Batto-Hokson · 5 real runs · 3 real job postings · Python · CrewAI · Ollama · pytest · October 2026

Scope note: the job postings are real (Brex and BILL, retrieved 2026-10-20 and summarized in my own words). The resume is a fictional candidate, so no real person's history is in the repo. Every output below is real and unedited.

EXECUTIVE SUMMARY

Free local models are good at rewording and bad at not making things up. Instead of asking a second AI to review the first one, this tool wraps a CrewAI crew (job description analyst, resume editor, cover letter writer) in plain Python checks. A draft that invents a number, adds a skill you can't defend, slips in a banned word or misses the cover letter's 7-to-9-sentence shape goes back to the model with the exact problems listed. After two failed retries it's saved marked FLAGGED instead of passing quietly.

I ran it five times with qwen2.5 (7B) on three real postings. All three resumes passed on the first attempt. Both cover letters were flagged after three attempts, one for inventing a "50%" improvement. Those runs also exposed three gaps in my own checks, which I fixed and covered with tests.

3 / 3	2 / 2	3	18
Resumes passed every check on attempt 1	Cover letters flagged, not passed off as done	Guard gaps found in real runs and fixed	Offline tests, no model needed

3 to 6 min

Per draft on 2 CPUs, $0 model cost

REAL RUNS

What a free 7B model did with real postings

Posting	Output	Result	What happened
Brex, Data Analyst II	Resume	**Passed, 1 try**	One edit: "exports 4 Tableau views" → "exports 4 dashboards", matching the posting
BILL, Risk Data Analyst	Resume	**Passed, 1 try**	Reordered and reworded a bullet; its change log misdescribed the edit
BILL, Business Development Rep	Resume	**Passed, 1 try**	No edits, and correctly didn't add Salesforce or ZoomInfo, which the candidate never listed
Brex, Data Analyst II	Cover letter	**Flagged, 3 tries**	10 sentences; invented "50%" (the real 3-days-to-1 cut is 67%)
BILL, Risk Data Analyst	Cover letter	**Flagged, 3 tries**	10 sentences, "leveraging", and "[Company Name]" left in

01 The guards kept an invented metric from passing

The final Brex cover letter claimed the candidate "improved lead follow-up efficiency by 50%". The resume says follow-up went from 3 days to 1 day. The model made up a number, and the number was wrong. The check compares every figure in the draft against the resume, job description and notes, so the letter failed and, after the last retry, was saved as FLAGGED.

```
Dear Hiring Manager,

I am excited to apply for the position of Data Analyst at your company. As an economics
graduate with a strong background in building data pipelines and transforming messy
data into actionable insights, I am eager to contribute to your team.
```

```
In my project "Retail Margin Analysis," I utilized Python, SQL, and Tableau to analyze
retail data over 84 weeks, uncovering a -1.5% gross margin on US bike sales. This
analysis helped identify key trends and inefficiencies, enabling more informed
decision-making.

At Bayside Motors, I improved lead follow-up efficiency by 50% through the development
of a call-back tracker in Google Sheets, reducing follow-up time from 3 days to 1 day.
This streamlined process directly contributed to a 14 vehicle sale average per month.

I am particularly interested in creating self-service data tools that empower business
teams to track their ...
```

WHY CODE, NOT ANOTHER AGENT

A second model reviewing the first shares its blind spots. A set difference between the numbers in the draft and the numbers in your sources doesn't.

02 Real runs found three holes in my own checks

What slipped through	Fix
A change log describing an edit the resume never got. It passed every check.	Each logged `New: "..."` edit must appear in the resume body
"leveraging" got past a banned list that contains "leverage"	Single banned words now match on their stem
"[Company Name]" placeholders in a finished letter	Bracketed placeholders are rejected

Each fix has a test built from the real output that exposed it. The original change-log failure is kept in `docs/found_in_testing/`.

STILL NOT CAUGHT

Change logs written in another format, generic openings like "I am excited to apply", and letters that never name the company. These are listed in the README's limits, not hidden.

03 A 7B model edits very conservatively

Across three resumes it made one real keyword edit. That's safe (nothing was fabricated), but it means the tool adds less value on a small model than on a stronger one. Switching to Claude is one line in `.env`; the checks don't change, so a stronger model just fails them less often.

DESIGN

How it works

```
Resume mode:   JD Analyst ──> Resume Editor ──> guards ──┬─ pass ──> output/<job>_resume.txt
                                                         └─ fail ──> retry with the exact
problems (max 2), then FLAGGED
Cover mode:    Hook Finder ──> Letter Writer ──> guards ──> output/<job>_cover.txt
```

Check	What it catches
New numbers	Any figure not in the resume (or, for letters, the job description and notes)
Blocked claims	Skills on your "can't defend this" list that the draft adds
New names	Tools or companies mid-sentence that appear in neither the resume nor the posting
Change log	Logged edits missing from the resume
Style	Em dashes, banned words and their variants, placeholders, code fences
Letter shape	7 to 9 sentences, no "I am writing...", no exclamation points

TECHNICAL REFERENCE

Setup and scope

Dimension	Value
Framework	CrewAI, sequential crews, one per mode
Model	Ollama qwen2.5:7b by default (free, local); Claude via one env var
Inputs	Real public postings (Brex, BILL) summarized with source links; a fictional resume
Testing	18 offline pytest tests with a scripted fake model; 5 real runs saved unedited
Tools	Python, CrewAI, Ollama, pytest